

Michael J. Pedersen : (908) 283-0318 : datacyclist@gmail.com : [My Site](#) : [LinkedIn](#)

Executive Summary

I am a [Google Cloud Certified Professional Data Engineer](#). I have 9 years of experience doing data engineering. I've scaled up a company from receiving 4T of new data/day through to its current 40T/day of new data. We have grown from about 50 servers to 450 servers for our big data architecture, as well as holding over 2P of data we are storing and actively using.

Relevant Technical Skills

- **Big Data:** [Google Big Query](#), [HDFS](#), [Hive](#), [YARN](#), [Alluxio](#), [Impala](#), [Trino](#), [Kafka](#), [Kubernetes](#), [Dagster](#), [Helm](#)
- **Database Servers:** [MySQL](#), [PostgreSQL](#), [Microsoft SQL Server](#)
- **Database Skills:** [PostgreSQL](#) Database Administration, Relational Schema Design, Structured Query Language (SQL)
- **Programming and Scripting Languages:** [Bash](#), C/C++, [Java](#), Javascript, [Perl](#), [PHP](#), [Python](#), [Terraform](#)
- **Programming Skills:** [Docker](#), [Jenkins](#), [Jira](#), [IntelliJ IDEA](#), Object-Oriented Design, Object-Oriented Programming, Refactoring
- **Software Configuration Management Tools:** [Git](#), [GitHub](#), [GitHub Actions](#), [Mercurial](#), [Subversion](#)

Relevant Job History

EvolutionIQ - Senior Software Engineer

New York City, NY - Feb 2024 - Current

- Created infrastructure to provide ephemeral environments, allowing testing of code branches before sending those code changes to the company at large.
- Orchestrated [Google Cloud](#) infrastructure with [Terraform](#), managing [GKE](#), [BigQuery](#), and developer virtual machines for scalability and efficiency.
- Redesigned EvolutionIQ's CI/CD pipeline, enhancing efficiency in development workflows.
- Fixed critical bugs in locally developed [GitHub Actions](#), ensuring smooth and reliable environment promotion.
- Built a template repository to guide EvolutionIQ teams in using [GitHub Actions](#) for deployment, as well as setting up data pipelines with [Dagster](#).
- Developed a Python-based CI/CD library to manage complex actions beyond [Bash](#)'s capabilities.
- Implemented [nektos/act](#) to enable developers to test build and deploy actions locally before committing to [GitHub](#).
- Refined [data retention policies](#) to guarantee compliance and ensure critical data availability for legal requirements.
- Implemented SLA monitoring with [Grafana](#), [SendGrid](#), and [PagerDuty](#) to proactively detect and alert customers and engineers about overnight job failures, ensuring swift issue resolution.
- On-call support on a rotating basis.

Pulsepoint - Data Engineer and Director of Infrastructure for Data

New York City, NY & Newark, NJ (Telecommute) - Mar 2015 - Nov 2023

Michael J. Pedersen : [\(908\) 283-0318](tel:9082830318) : datacyclist@gmail.com : [My Site](#) : [LinkedIn](#)

Director of Infrastructure for Data, May 2018 - Nov 2023

- Architected data streaming that manages 40T of data/day.
- Lead maintainer for ETL pipelines, encompassing over 250 transformations.
- Established new data centers in Europe and in Virginia.
- Migrated data center, moving processing of data pipelines to new data center.
- Guided the team through splitting our ETL pipelines into multiple repositories.
- Organized the migration of ETL pipelines from Python 2 to Python 3.
- Replaced [Vertica](#) with [Trino](#).
- Tested new tools for suitability, including [MariaDB](#), [Clickhouse](#), and [Kudu](#).
- Changed hardware profiles for [Hadoop](#) to remove storage and compute colocation.
- Passed annual HIPAA training for data protection.
- Reported on system wide data latency using [ElasticSearch](#), [Kibana](#), and [Grafana](#).

Data Engineer, Mar 2015 - May 2018

- Built tool to graphically show the ETL pipelines.
- Created ELT jobs to ingest third party data to make it available internally.
- Troubleshooting of issues with [Hadoop](#), [Kafka](#), [SQL Server](#), and [Kubernetes](#).
- Production maintenance of data pipelines, including after hours support.
- Installed and configured multiple [Hadoop](#) clusters.
- Implemented data duplication between two [Hadoop](#) clusters.
- Upgraded [Hadoop](#) clusters with minimal downtime.
- Tested [Cassandra](#) as a potential reporting database.
- Transitioned ETL pipeline from cronjobs to [Mesos](#) and then into [Kubernetes](#).
- Converted [Sqoop](#) jobs to use [FreeBCP](#) instead.
- Collaborated with other teams to help them use the systems to find the data they need.
- Optimized the performance and reliability of [Hadoop](#), ensuring high availability.
- Worked with other teams to define and then implement needed features for our internal ETL pipeline framework.
- Added over a hundred automated tests to our ETL pipeline.
- Performed root cause analysis on ETL and cluster level failures.
- Managed data backfill issues whenever they arose.

Weight Watchers - Systems Engineering Lead

New York City, NY - Nov 2014 - Feb 2015

- Developed lightweight monitoring tool for use within my group.
- Configured [Vormetric](#) products to ensure [HIPAA](#) compliance for customer data.
- Worked to transfer from [Rackspace Cloud](#) to [Openstack](#) based private cloud.

OrcaTec, LLC - Developer

Atlanta, GA (Telecommute) - Jun 2012 - Oct 2014

- Reduced multi-hour [SQLAlchemy](#) bulk database jobs to minutes.
- Added holds and matters framework, allowing customers to state that documents belong to specific cases and should not be deleted while the cases are ongoing.
- Wrote [Python](#) framework to manage long running background jobs.
- Debugged and resolved memory issues that were causing systems to shut down.

Choopa.com - Developer

Sayreville, NJ - Jan 2012 - May 2012

Michael J. Pedersen : [\(908\) 283-0318](tel:9082830318) : datacyclist@gmail.com : [My Site](#) : [LinkedIn](#)

- Developed library to manage [OpenStack](#) nodes, and gather billing information.
- Built [Nagios](#) configuration file generator for in-house web interface for [Nagios](#).
- Configured [Bacula](#) backup system as replacement for custom backup scripts.
- Reconfigured [Nagios](#) monitoring, reducing full check from 8 hours to 2 minutes.
- Refactored in-house [Nagios](#) web interface. This reduced the workload from six files down to one when adding new checks.
- Several smaller bug fixes and features throughout the internal code base.

Education

Bachelor of Science in Computer Science, 2000
East Stroudsburg University, East Stroudsburg, Pennsylvania

Professional Certificates

[Google Cloud Certified Professional Data Engineer](#) (Jan 2024)
Online Course - Google